
Amazon Redshift

Getting Started Guide

API Version 2012-12-01



Amazon Redshift: Getting Started Guide

Copyright © 2018 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

Getting Started	1
Step 1: Set Up Prerequisites	1
Sign Up for AWS	2
Determine Firewall Rules	2
Step 2: Create an IAM Role	2
To Create an IAM Role for Amazon Redshift	2
Step 3: Launch a Cluster	3
To Launch an Amazon Redshift Cluster	3
Step 4: Authorize Cluster Access	5
To Configure the VPC Security Group (EC2-VPC Platform)	5
To Configure the Amazon Redshift Security Group	6
Step 5: Connect to the Cluster	7
Querying a Database Using the Query Editor	7
Querying a Database Using a SQL Client	10
Step 6: Load Sample Data	13
Step 7: Find Additional Resources and Reset Your Environment	17
Where Do I Go From Here?	18
Document History	21

Getting Started with Amazon Redshift

Welcome to the *Amazon Redshift Getting Started Guide*. Amazon Redshift is a fully managed, petabyte-scale data warehouse service in the cloud. An Amazon Redshift data warehouse is a collection of computing resources called *nodes*, which are organized into a group called a *cluster*. Each cluster runs an Amazon Redshift engine and contains one or more databases.

If you are a first-time user of Amazon Redshift, we recommend that you begin by reading the following sections:

- [Amazon Redshift Management Overview](#) – This topic provides an overview of Amazon Redshift.
- [Service Highlights and Pricing](#) – This product detail page provides the Amazon Redshift value proposition, service highlights, and pricing.
- [Amazon Redshift Getting Started \(this guide\)](#) – This guide provides a tutorial of using Amazon Redshift to create a sample cluster and work with sample data.

If you are building a proof-of-concept solution with Amazon Redshift, we recommend that you read [Building a Proof of Concept for Amazon Redshift](#).

This guide is a tutorial designed to walk you through the process of creating a sample Amazon Redshift cluster. You can use this sample cluster to evaluate the Amazon Redshift service. In this tutorial, you perform the following steps:

- [Step 1: Set Up Prerequisites \(p. 1\)](#)
- [Step 2: Create an IAM Role \(p. 2\)](#)
- [Step 3: Launch a Sample Amazon Redshift Cluster \(p. 3\)](#)
- [Step 4: Authorize Access to the Cluster \(p. 5\)](#)
- [Step 5: Connect to the Sample Cluster and Run Queries \(p. 7\)](#)
- [Step 6: Load Sample Data from Amazon S3 \(p. 13\)](#)
- [Step 7: Find Additional Resources and Reset Your Environment \(p. 17\)](#)

After you complete this tutorial, you can find more information about Amazon Redshift and next steps in [Where Do I Go From Here? \(p. 18\)](#)

Important

The sample cluster that you create will be running in a live environment. The on-demand rate is \$0.25 per hour for using the sample cluster that is designed in this tutorial until you delete it. For more pricing information, go to [the Amazon Redshift pricing page](#). If you have questions or get stuck, you can reach out to the Amazon Redshift team by posting on our [Discussion Forum](#).

This tutorial is not meant for production environments, and does not discuss options in depth. After you complete the steps in this tutorial, you can use the [Additional Resources \(p. 18\)](#) section to locate more in-depth information to plan, deploy, and maintain your clusters, and to work with the data in your data warehouse.

Step 1: Set Up Prerequisites

Before you begin setting up an Amazon Redshift cluster, make sure that you complete the following prerequisites in this section:

- [Sign Up for AWS](#) (p. 2)
- [Determine Firewall Rules](#) (p. 2)

Sign Up for AWS

If you don't already have an AWS account, you must sign up for one. If you already have an account, you can skip this prerequisite and use your existing account.

1. Open <https://aws.amazon.com/>, and then choose **Create an AWS Account**.

Note

If you previously signed in to the AWS Management Console using AWS account root user credentials, choose **Sign in to a different account**. If you previously signed in to the console using IAM credentials, choose **Sign-in using root account credentials**. Then choose **Create a new AWS account**.

2. Follow the online instructions.

Part of the sign-up procedure involves receiving a phone call and entering a verification code using the phone keypad.

Determine Firewall Rules

As part of this tutorial, you will specify a port when you launch your Amazon Redshift cluster. You will also create an inbound ingress rule in a security group to allow access through the port to your cluster.

If your client computer is behind a firewall, you need to know an open port that you can use so you can connect to the cluster from a SQL client tool and run queries. If you do not know this, you should work with someone who understands your network firewall rules to determine an open port in your firewall. Though Amazon Redshift uses port 5439 by default, the connection will not work if that port is not open in your firewall. Because you cannot change the port number for your Amazon Redshift cluster after it is created, make sure that you specify an open port that will work in your environment during the launch process.

Step 2: Create an IAM Role

For any operation that accesses data on another AWS resource, such as using a COPY command to load data from Amazon S3, your cluster needs permission to access the resource and the data on the resource on your behalf. You provide those permissions by using AWS Identity and Access Management, either through an IAM role that is attached to your cluster or by providing the AWS access key for an IAM user that has the necessary permissions.

To best protect your sensitive data and safeguard your AWS access credentials, we recommend creating an IAM role and attaching it to your cluster. For more information about providing access permissions, see [Permissions to Access Other AWS Resources](#).

In this step, you will create a new IAM role that enables Amazon Redshift to load data from Amazon S3 buckets. In the next step, you will attach the role to your cluster.

To Create an IAM Role for Amazon Redshift

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.

2. In the left navigation pane, choose **Roles**.
3. Choose **Create role**
4. In the **AWS Service** group, choose **Redshift**.
5. Under **Select your use case**, choose **Redshift - Customizable** then choose **Next: Permissions**.
6. On the **Attach permissions policies** page, choose **AmazonS3ReadOnlyAccess**, and then choose **Next: Review**.
7. For **Role name**, type a name for your role. For this tutorial, type `myRedshiftRole`.
8. Review the information, and then choose **Create Role**.
9. Choose the role name of the role you just created.
10. Copy the **Role ARN** to your clipboard—this value is the Amazon Resource Name (ARN) for the role that you just created. You will use that value when you use the COPY command to load data in [Step 6: Load Sample Data from Amazon S3 \(p. 13\)](#).

Now that you have created the new role, your next step is to attach it to your cluster. You can attach the role when you launch a new cluster or you can attach it to an existing cluster. In the next step, you'll attach the role to a new cluster.

Step 3: Launch a Sample Amazon Redshift Cluster

Now that you have the prerequisites completed, you can launch your Amazon Redshift cluster.

Important

The cluster that you are about to launch will be live (and not running in a sandbox). You will incur the standard Amazon Redshift usage fees for the cluster until you delete it. If you complete the tutorial described here in one sitting and delete the cluster when you are finished, the total charges will be minimal.


To Launch an Amazon Redshift Cluster

1. Sign in to the AWS Management Console and open the Amazon Redshift console at <https://console.aws.amazon.com/redshift/>.

Important

If you use IAM user credentials, ensure that the user has the necessary permissions to perform the cluster operations. For more information, go to [Controlling Access to IAM Users](#) in the *Amazon Redshift Cluster Management Guide*.

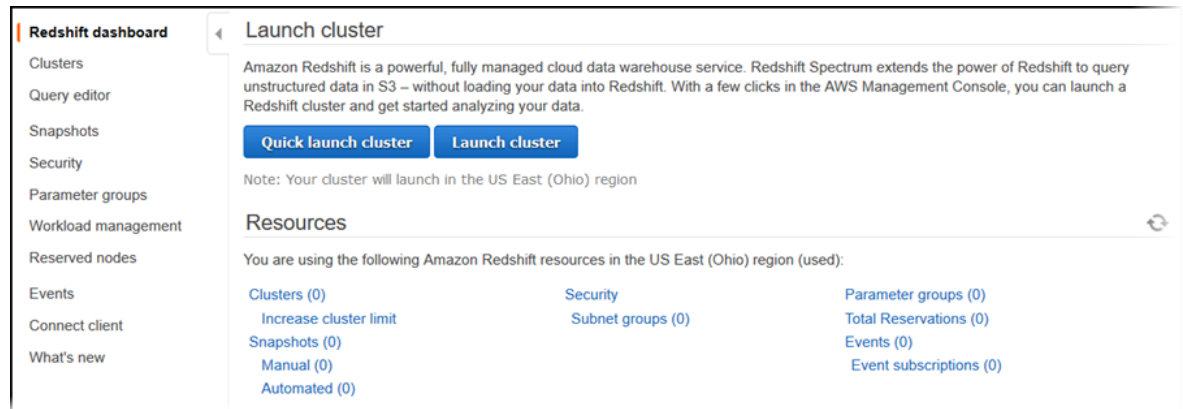
2. In the main menu, select the region in which you want to create the cluster. For the purposes of this tutorial, select **US West (Oregon)**.



3. On the Amazon Redshift Dashboard, choose **Quick launch cluster**.

The Amazon Redshift Dashboard looks similar to the following.


Amazon Redshift Getting Started Guide To Launch an Amazon Redshift Cluster



4. On the Cluster specifications page, enter the following values and then choose **Launch cluster**:
- **Node type**: Choose **dc2.large**.
 - **Number of compute nodes**: Keep the default value of **2**.
 - **Master user name**: Keep the default value of **awsuser**.
 - **Master user password** and **Confirm password**: Enter a password for the master user account.
 - **Database port**: Accept the default value of **5439**.
 - **Available IAM roles**: Choose **myRedshiftRole**.

The screenshot shows the 'Launch your Amazon Redshift cluster - Quick launch' page. It has a heading 'Launch your Amazon Redshift cluster - Quick launch' and a link to 'Switch to advanced settings'. Below this is the 'Cluster specifications' section. It contains several input fields: 'Node type*' with a dropdown menu showing 'dc2.large' and additional info like 'Storage type: SSD', 'Storage: 0.16 TB/node', and 'Compute optimized'; 'Number of compute nodes*' with a spinner set to '2' and a calculation 'x 0.16 TB/node = 0.32 TB storage available'; 'Cluster identifier*' with a text box containing 'redshift-cluster-1'; 'Master user name*' with a text box containing 'awsuser'; 'Master user password*' and 'Confirm password*' with masked password fields; 'Database port*' with a text box containing '5439'; and 'Available IAM roles' with a dropdown menu showing 'Choose a role' and a radio button selected for 'myRedshiftRole'. At the bottom right are 'Cancel' and 'Launch cluster' buttons.

5. A confirmation page appears and the cluster will take a few minutes to finish. Choose **Close** to return to the list of clusters.



- On the Clusters page, choose the cluster that you just launched and review the **Cluster Status** information. Make sure that the **Cluster Status** is **available** and the **Database Health** is **healthy** before you try to connect to the database later in this tutorial.

The screenshot shows the Amazon Redshift Clusters console. At the top, there are buttons for "Quick launch cluster" and "Launch cluster", followed by dropdown menus for "Cluster", "Database", and "Backup", and buttons for "Manage Tags" and "Manage IAM roles". Below this is a table with columns: "Cluster", "Cluster Status", "DB Health", "Release Status", and "In Maintenance". There are two rows of cluster data.

	Cluster	Cluster Status	DB Health	Release Status	In Maintenance
<input type="checkbox"/>	[redacted]	available	healthy	Up to date	no
<input type="checkbox"/>	redshift-cluster-1	available	healthy	Up to date	no

Step 4: Authorize Access to the Cluster

In the previous step, you launched your Amazon Redshift cluster. Before you can connect to the cluster, you need to configure a security group to authorize access:


- If you launched your cluster in the EC2-VPC platform, follow the steps in [To Configure the VPC Security Group \(EC2-VPC Platform\) \(p. 5\)](#).
- If you launched your cluster in the EC2-Classic platform, follow the steps in [To Configure the Amazon Redshift Security Group \(p. 6\)](#).

Note


You only need to configure one of these two types of security groups. Follow the steps that correspond to the platform in which you launched your cluster.

To Configure the VPC Security Group (EC2-VPC Platform)

- In the Amazon Redshift console, in the navigation pane, choose **Clusters**.
- Choose **examplecluster** to open it, and make sure you are on the **Configuration** tab.
- Under **Cluster Properties**, for **VPC Security Groups**, choose your security group.




- After your security group opens in the Amazon EC2 console, choose the **Inbound** tab.



5. Choose **Edit**, and enter the following, then choose **Save**:
 - **Type**: Custom TCP Rule.
 - **Protocol**: TCP.
 - **Port Range**: type the same port number that you used when you launched the cluster. The default port for Amazon Redshift is 5439, but your port might be different.
 - **Source**: select **Custom IP**, then type 0.0.0.0/0.


Important

Using 0.0.0.0/0 is not recommended for anything other than demonstration purposes because it allows access from any computer on the internet. In a real environment, you would create inbound rules based on your own network settings.




To Configure the Amazon Redshift Security Group

1. In the Amazon Redshift console, in the navigation pane, choose **Clusters**.
2. Choose `examplecluster` to open it, and make sure you are on the **Configuration** tab.
3. Under **Cluster Properties**, for **Cluster Security Groups**, choose **default** to open the default security group.



4. On the **Security Groups** tab, in the cluster security group list, choose the cluster security group whose rules you want to manage.
5. On the **Security Group Connections** tab, choose **Add Connection Type**.

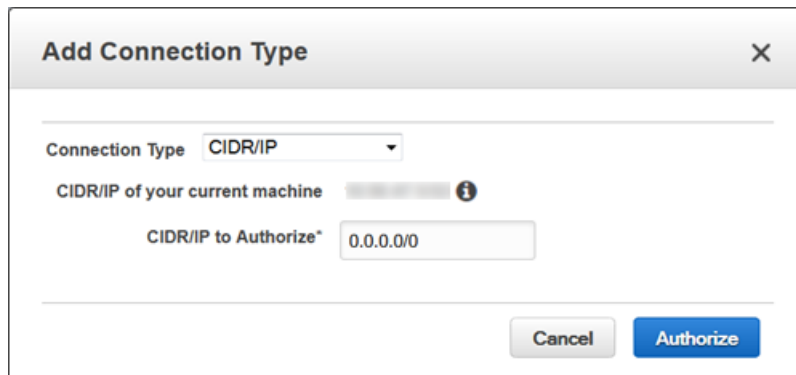


6. In the **Connection Type** box, choose **CIDR/IP**.

In **CIDR/IP to Authorize**, type `0.0.0.0/0` and choose **Authorize**.

Important

Using `0.0.0.0/0` is not recommended for anything other than demonstration purposes because it allows access from any computer on the Internet. In a real environment, you would create inbound rules based on your own network settings.



Step 5: Connect to the Sample Cluster and Run Queries

To query databases hosted by your Amazon Redshift cluster, you have two options:

- Connect to your cluster and run queries on the AWS Management Console with the Query Editor.

If you use the Query Editor, you don't have to download and set up a SQL client application.

- Connect to your cluster through a SQL client tool, such as SQL Workbench/j.

Topics

- [Querying a Database Using the Query Editor \(p. 7\)](#)
- [Querying a Database Using a SQL Client \(p. 10\)](#)

Querying a Database Using the Query Editor

Using the Query Editor is the easiest way to run queries on databases hosted by your Amazon Redshift cluster. After creating your cluster, you can immediately run queries by using the Query Editor on the Amazon Redshift console.

The following cluster node types support the Query Editor:

- DC1.8xlarge

- DC2.large
- DC2.8xlarge
- DS2.8xlarge

Using the Query Editor, you can do the following:

- Run single SQL statement queries.
- Download result sets as large as 100 MB to a comma-separated value (CSV) file.
- Save queries for reuse. You can't save queries in the EU (Paris) Region or the Asia Pacific (Osaka-Local) Region.
- View query execution details for user-defined tables.

Query Editor Considerations

Be aware of the following considerations when you use the Query Editor:

- Up to 50 users can connect to a cluster with the Query Editor at the same time.
- Up to 500 users can connect to a cluster simultaneously. This total includes the users connecting through the Query Editor.
- Up to 50 workload management (WLM) query slots can be active at the same time. For more information about query slots, see [Implementing Workload Management](#).
- Query Editor only runs short queries that can complete within two minutes.
- Query result sets are paginated with 100 rows per page.
- You can't use the Query Editor with Enhanced VPC routing.
- You can't use transactions in the Query Editor. For more information about transactions, see [BEGIN](#) in the *Amazon Redshift Database Developer Guide*.

Enabling Access to the Query Editor

To access the Query Editor, you need permission. To enable access, attach the `AmazonRedshiftQueryEditor` and `AmazonRedshiftReadOnlyAccess` policies for AWS Identity and Access Management (IAM) to the AWS account that you use to access your cluster.

If you have already created an IAM user to access Amazon Redshift, you can attach the `AmazonRedshiftQueryEditor` and `AmazonRedshiftReadOnlyAccess` policies to that user. If you haven't created an IAM user yet, create one and attach the policies to the IAM user.

To attach the required IAM policies for the Query Editor

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. Choose **Users**.
3. Choose the user that needs access to the Query Editor.
4. Choose **Add permissions**.
5. Choose **Attach existing policies directly**.
6. For **Policy names**, choose **AmazonRedshiftQueryEditor** and **AmazonRedshiftReadOnlyAccess**.
7. Choose **Next: Review**.
8. Choose **Add permissions**.

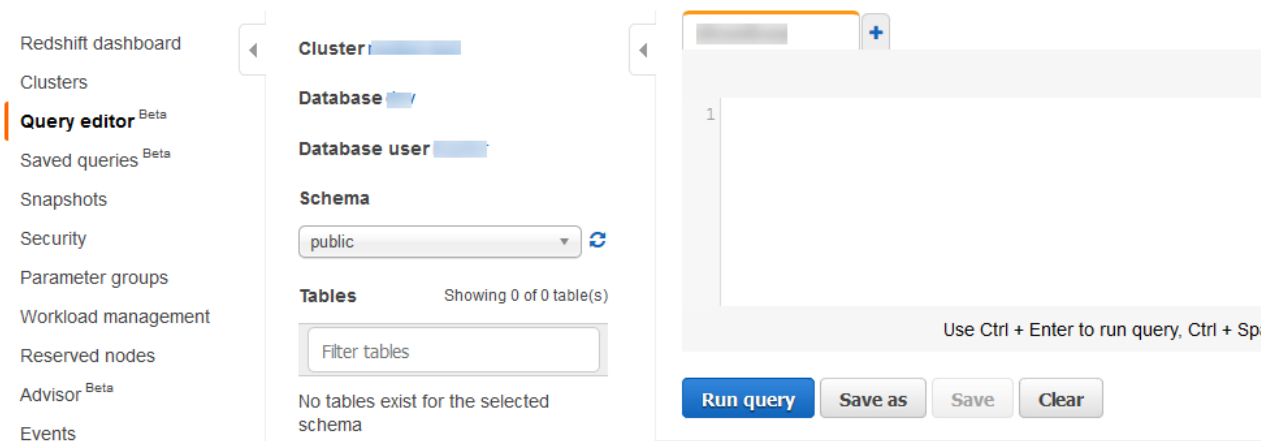
Using the Query Editor

In the following example, you use the Query Editor to perform the following tasks:

- Run SQL commands.
- View query execution details.
- Save a query.
- Download a query result set.

To use the Query Editor

1. Sign in to the AWS Management Console and open the Amazon Redshift console at <https://console.aws.amazon.com/redshift/>.
2. In the navigation pane, choose **Query Editor**.



3. For **Schema**, choose **public** to create a new table based on that schema.
4. Enter the following in the Query Editor window and choose **Run query** to create a new table.

```
create table shoes(  
  shoetype varchar (10),  
  color varchar(10));
```

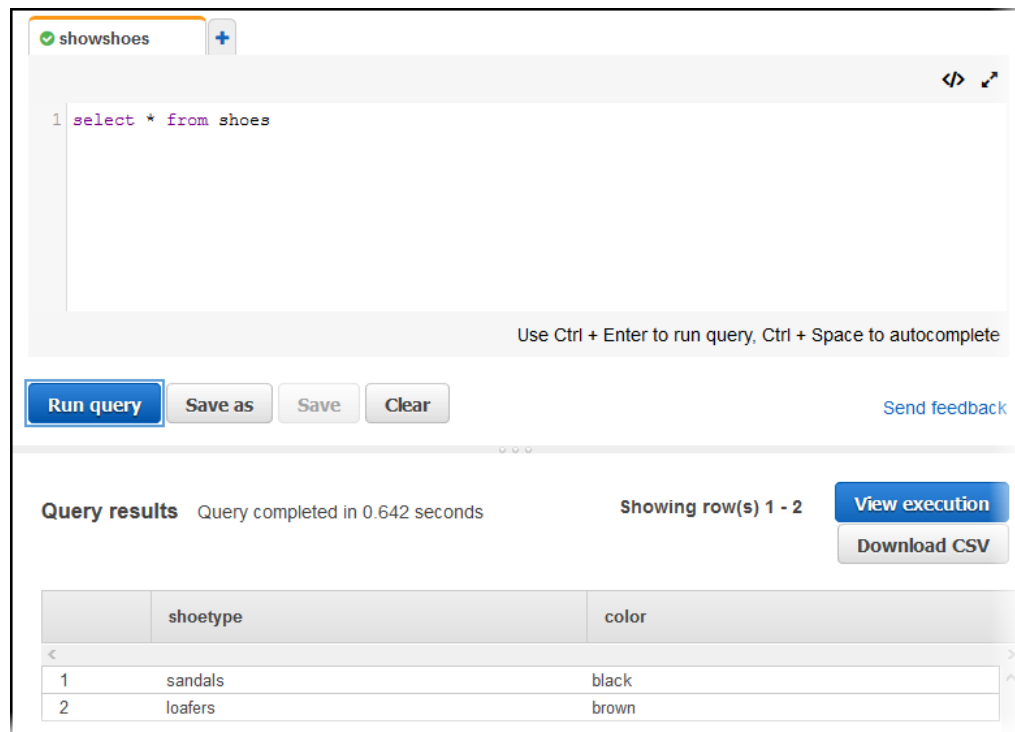
5. Choose **Clear**.
6. Enter the following command in the Query Editor window and choose **Run query** to add rows to the table.

```
insert into shoes values  
( 'loafers', 'brown'),  
( 'sandals', 'black');
```

7. Choose **Clear**.
8. Enter the following command in the Query Editor window and choose **Run query** to query the new table.

```
select * from shoes;
```

You should see the following results.



Querying a Database Using a SQL Client

Now you will connect to your cluster by using a SQL client tool and run a simple query to test the connection. You can use most SQL client tools that are compatible with PostgreSQL. For this tutorial, you'll use the SQL Workbench/J client that you installed in the prerequisites section of this tutorial. Complete this section by performing the following steps:

- [Install SQL Client Drivers and Tools \(p. 10\)](#)
- [To Get Your Connection String \(p. 11\)](#)
- [To Connect from SQL Workbench/J to Your Cluster \(p. 11\)](#)

After you complete this step, you can determine whether you want to load sample data from Amazon S3 in [Step 6: Load Sample Data from Amazon S3 \(p. 13\)](#) or find more information about Amazon Redshift and reset your environment at [Where Do I Go From Here? \(p. 18\)](#).

Install SQL Client Drivers and Tools

You can use most SQL client tools with Amazon Redshift JDBC or ODBC drivers to connect to an Amazon Redshift cluster. In this tutorial, we show you how to connect using SQL Workbench/J, a free, DBMS-independent, cross-platform SQL query tool. If you plan to use SQL Workbench/J to complete this tutorial, follow the steps below to get set up with the Amazon Redshift JDBC driver and SQL Workbench/J. For more complete instructions for installing SQL Workbench/J, go to [Setting Up the SQL Workbench/J Client](#) in the *Amazon Redshift Cluster Management Guide*. If you use an Amazon EC2 instance as your client computer, you will need to install SQL Workbench/J and the required drivers on the instance.

Note

You must install any third-party database tools that you want to use with your clusters; Amazon Redshift does not provide or install any third-party tools or libraries.

To Install SQL Workbench/J on Your Client Computer

1. Review the [SQL Workbench/J software license](#).
2. Go to the [SQL Workbench/J website](#) and download the appropriate package for your operating system.
3. Go to the [Installing and starting SQL Workbench/J page](#) and install SQL Workbench/J.

Important

Note the Java runtime version prerequisites for SQL Workbench/J and ensure you are using that version, otherwise, this client application will not run.

4. Go to [Configure a JDBC Connection](#) and download an Amazon Redshift JDBC driver to enable SQL Workbench/J to connect to your cluster.

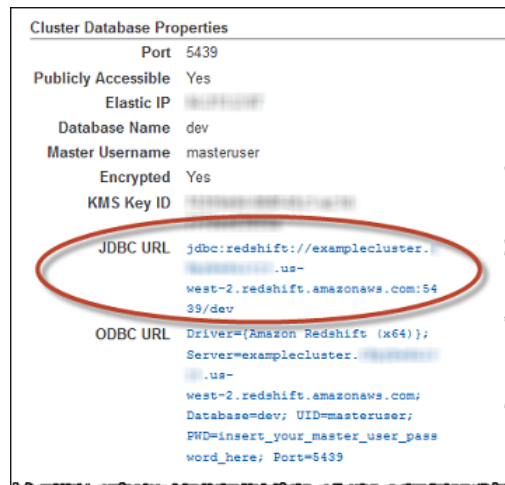
For more information about using the Amazon Redshift JDBC or ODBC drivers, see [Configuring Connections in Amazon Redshift](#).

To Get Your Connection String

1. In the Amazon Redshift console, in the navigation pane, choose **Clusters**.
2. Choose `examplecluster` to open it, and make sure you are on the **Configuration** tab.
3. On the **Configuration** tab, under **Cluster Database Properties**, copy the JDBC URL of the cluster.

Note

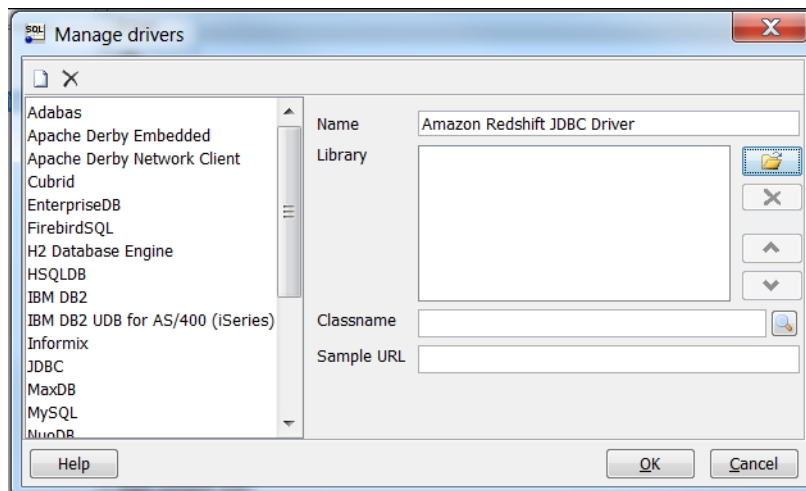
The endpoint for your cluster is not available until the cluster is created and in the available state.




To Connect from SQL Workbench/J to Your Cluster

This step assumes you installed SQL Workbench/J in [Step 1: Set Up Prerequisites \(p. 1\)](#).

1. Open SQL Workbench/J.
2. Choose **File**, and then choose **Connect window**.
3. Choose **Create a new connection profile**.
4. In the **New profile** text box, type a name for the profile.
5. Choose **Manage Drivers**. The **Manage Drivers** dialog opens.
6. Choose the **Create a new entry** button. In the **Name** text box, type a name for the driver.




Choose the folder icon next to the **Library** box, navigate to the location of the driver, select it, and then choose **Open**.




If the **Please select one driver** dialog box displays, select **com.amazon.redshift.jdbc4.Driver** or **com.amazon.redshift.jdbc41.Driver** and choose **OK**. SQL Workbench/J automatically completes the **Classname** box. Leave the **Sample URL** box blank, and then choose **OK**.

7. In the **Driver** box, choose the driver you just added.
8. In **URL**, copy the JDBC URL from the Amazon Redshift console and paste it here.
9. In **Username**, type *masteruser*.
10. In **Password**, type the password associated with the master user account.
11. Choose the **Autocommit** box.
12. Choose the **Save profile list** icon, as shown below:



13. Choose **OK**.



14. Enter the following command in the query window and choose **SQL, Execute Current** to add rows to the table.

```
create table shoes(  
  shoetype varchar (10),  
  color varchar(10));
```

15. Run the following command to add rows to the table.

```
insert into shoes values  
( 'loafers', 'brown'),  
( 'sandals', 'black');
```

16. Run the following command to query the new table.

```
select * from shoes;
```

Step 6: Load Sample Data from Amazon S3

At this point you have a database called dev and you are connected to it. Now you will create some tables in the database, upload data to the tables, and try a query. For your convenience, the sample data you will load is available in an Amazon S3 bucket.

Note

If you're using a SQL client tool, ensure that your SQL client is connected to the cluster.

After you complete this step, you can find more information about Amazon Redshift and reset your environment at [Where Do I Go From Here? \(p. 18\)](#).

1. Create tables.

Copy and execute the following create table statements to create tables in the dev database. For more information about the syntax, go to [CREATE TABLE](#) in the *Amazon Redshift Database Developer Guide*.

```
create table users(
  userid integer not null distkey sortkey,
  username char(8),
  firstname varchar(30),
  lastname varchar(30),
  city varchar(30),
  state char(2),
  email varchar(100),
  phone char(14),
  likesports boolean,
  liketheatre boolean,
  likeconcerts boolean,
  likejazz boolean,
  likeclassical boolean,
  likeopera boolean,
  likerock boolean,
  likevegas boolean,
  likebroadway boolean,
  likemusicals boolean);

create table venue(
  venueid smallint not null distkey sortkey,
  venuename varchar(100),
  venuecity varchar(30),
  venuestate char(2),
  venueseats integer);

create table category(
  catid smallint not null distkey sortkey,
  catgroup varchar(10),
  catname varchar(10),
  catdesc varchar(50));

create table date(
  dateid smallint not null distkey sortkey,
  caldate date not null,
  day character(3) not null,
  week smallint not null,
  month character(5) not null,
  qtr character(5) not null,
  year smallint not null,
  holiday boolean default('N'));

create table event(
  eventid integer not null distkey,
  venueid smallint not null,
  catid smallint not null,
  dateid smallint not null sortkey,
  eventname varchar(200),
  starttime timestamp);

create table listing(
  listid integer not null distkey,
  sellerid integer not null,
  eventid integer not null,
  dateid smallint not null sortkey,
  numtickets smallint not null,
  priceperticket decimal(8,2),
```

```
totalprice decimal(8,2),
listtime timestamp);

create table sales(
  salesid integer not null,
  listid integer not null distkey,
  sellerid integer not null,
  buyerid integer not null,
  eventid integer not null,
  dateid smallint not null sortkey,
  qty sold smallint not null,
  pricepaid decimal(8,2),
  commission decimal(8,2),
  saletime timestamp);
```

2. Load sample data from Amazon S3 by using the COPY command.

Note

We recommend using the COPY command to load large datasets into Amazon Redshift from Amazon S3 or DynamoDB. For more information about COPY syntax, see [COPY](#) in the *Amazon Redshift Database Developer Guide*.

The sample data for this tutorial is provided in an Amazon S3 bucket that is owned by Amazon Redshift. The bucket permissions are configured to allow all authenticated AWS users read access to the sample data files.

To load the sample data, you must provide authentication for your cluster to access Amazon S3 on your behalf. You can provide either role-based authentication or key-based authentication. We recommend using role-based authentication. For more information about both types of authentication, see [CREDENTIALS](#) in the Amazon Redshift Database Developer Guide.

For this step, you will provide authentication by referencing the IAM role you created and then attached to your cluster in previous steps.

Note

If you don't have proper permissions to access Amazon S3, you receive the following error message when running the COPY command: `S3ServiceException: Access Denied`.

The COPY commands include a placeholder for the IAM role ARN, as shown in the following example.

```
copy users from 's3://awssampleduswest2/ticket/allusers_pipe.txt'
credentials 'aws_iam_role=<iam-role-arn>'
delimiter '|' region 'us-west-2';
```

To authorize access using an IAM role, replace `<iam-role-arn>` in the CREDENTIALS parameter string with the role ARN for the IAM role you created in [Step 2: Create an IAM Role \(p. 2\)](#).

Your COPY command will look similar to the following example.

```
copy users from 's3://awssampleduswest2/ticket/allusers_pipe.txt'
credentials 'aws_iam_role=arn:aws:iam::123456789012:role/myRedshiftRole'
delimiter '|' region 'us-west-2';
```

To load the sample data, replace `<iam-role-arn>` in the following COPY commands with your role ARN. Then run the commands in your SQL client tool.

```
copy users from 's3://awssampleduswest2/ticket/allusers_pipe.txt'
credentials 'aws_iam_role=<iam-role-arn>'
delimiter '|' region 'us-west-2';
```

```
copy venue from 's3://awssampleduswest2/ticket/venue_pipe.txt'
credentials 'aws_iam_role=<iam-role-arn>'
delimiter '|' region 'us-west-2';

copy category from 's3://awssampleduswest2/ticket/category_pipe.txt'
credentials 'aws_iam_role=<iam-role-arn>'
delimiter '|' region 'us-west-2';

copy date from 's3://awssampleduswest2/ticket/date2008_pipe.txt'
credentials 'aws_iam_role=<iam-role-arn>'
delimiter '|' region 'us-west-2';

copy event from 's3://awssampleduswest2/ticket/allevvents_pipe.txt'
credentials 'aws_iam_role=<iam-role-arn>'
delimiter '|' timeformat 'YYYY-MM-DD HH:MI:SS' region 'us-west-2';

copy listing from 's3://awssampleduswest2/ticket/listings_pipe.txt'
credentials 'aws_iam_role=<iam-role-arn>'
delimiter '|' region 'us-west-2';

copy sales from 's3://awssampleduswest2/ticket/sales_tab.txt'
credentials 'aws_iam_role=<iam-role-arn>'
delimiter '\t' timeformat 'MM/DD/YYYY HH:MI:SS' region 'us-west-2';
```

3. Now try the example queries. For more information, go to [SELECT](#) in the *Amazon Redshift Developer Guide*.

```
-- Get definition for the sales table.
SELECT *
FROM pg_table_def
WHERE tablename = 'sales';

-- Find total sales on a given calendar date.
SELECT sum(qtysold)
FROM   sales, date
WHERE  sales.dateid = date.dateid
AND    caldate = '2008-01-05';

-- Find top 10 buyers by quantity.
SELECT firstname, lastname, total_quantity
FROM   (SELECT buyerid, sum(qtysold) total_quantity
        FROM   sales
        GROUP BY buyerid
        ORDER BY total_quantity desc limit 10) Q, users
WHERE  Q.buyerid = userid
ORDER BY Q.total_quantity desc;

-- Find events in the 99.9 percentile in terms of all time gross sales.
SELECT eventname, total_price
FROM   (SELECT eventid, total_price, ntile(1000) over(order by total_price desc) as
        percentile
        FROM (SELECT eventid, sum(pricepaid) total_price
              FROM   sales
              GROUP BY eventid)) Q, event E
WHERE  Q.eventid = E.eventid
AND    percentile = 1
ORDER BY total_price desc;
```

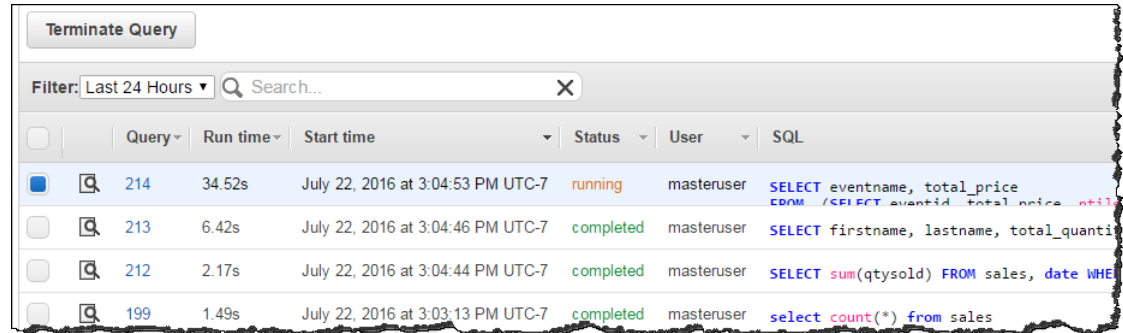
4. You can optionally go the Amazon Redshift console to review the queries you executed. The **Queries** tab shows a list of queries that you executed over a time period you specify. By default, the console displays queries that have executed in the last 24 hours, including currently executing queries.

Amazon Redshift Getting Started Guide

Step 7: Find Additional Resources and Reset Your Environment

- Sign in to the AWS Management Console and open the Amazon Redshift console at <https://console.aws.amazon.com/redshift/>.
- In the cluster list in the right pane, choose `examplecluster`.
- Choose the **Queries** tab.

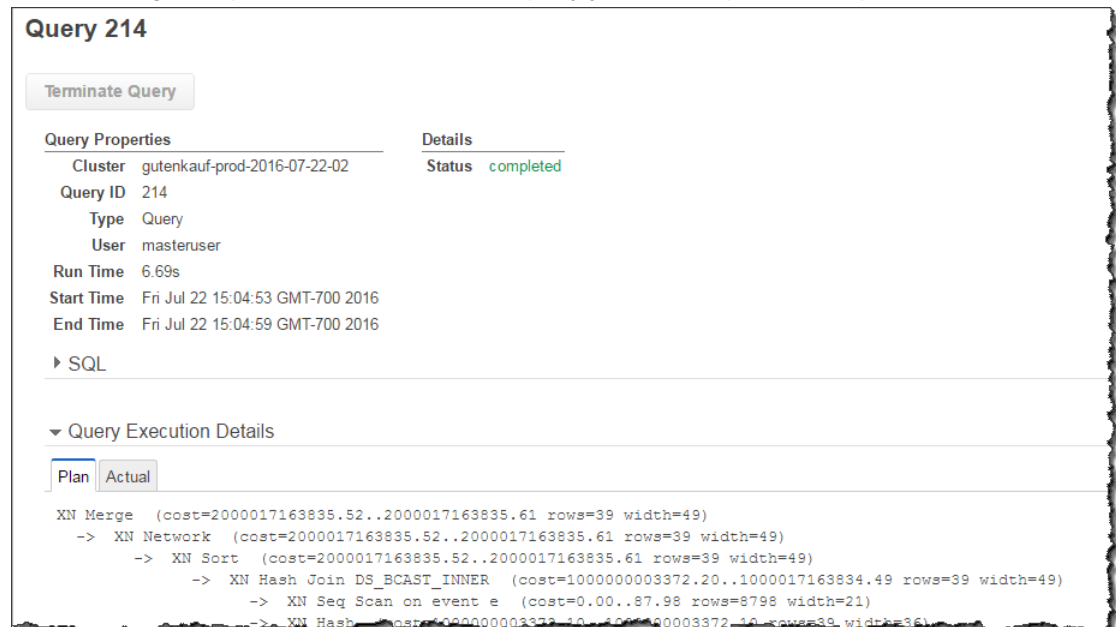
The console displays list of queries you executed as shown in the example below.



Query	Run time	Start time	Status	User	SQL
214	34.52s	July 22, 2016 at 3:04:53 PM UTC-7	running	masteruser	SELECT eventname, total_price FROM ...
213	6.42s	July 22, 2016 at 3:04:46 PM UTC-7	completed	masteruser	SELECT firstname, lastname, total_quant...
212	2.17s	July 22, 2016 at 3:04:44 PM UTC-7	completed	masteruser	SELECT sum(qtysold) FROM sales, date WHE...
199	1.49s	July 22, 2016 at 3:03:13 PM UTC-7	completed	masteruser	select count(*) from sales

- To view more information about a query, choose the query ID link in the **Query** column or choose the magnifying glass icon.

The following example shows the details of a query you ran in a previous step.



Query 214

Terminate Query

Query Properties	Details
Cluster: gutenkauf-prod-2016-07-22-02	Status: completed
Query ID: 214	
Type: Query	
User: masteruser	
Run Time: 6.69s	
Start Time: Fri Jul 22 15:04:53 GMT-700 2016	
End Time: Fri Jul 22 15:04:59 GMT-700 2016	

SQL

Query Execution Details

Plan Actual

```
XN Merge (cost=2000017163835.52..2000017163835.61 rows=39 width=49)
-> XN Network (cost=2000017163835.52..2000017163835.61 rows=39 width=49)
-> XN Sort (cost=2000017163835.52..2000017163835.61 rows=39 width=49)
-> XN Hash Join DS_BCAST_INNER (cost=1000000003372.20..1000017163834.49 rows=39 width=49)
-> XN Seq Scan on event e (cost=0.00..87.98 rows=8798 width=21)
-> XN Hash (cost=1000000003372.20..1000017163834.49 rows=39 width=36)
```

Step 7: Find Additional Resources and Reset Your Environment

When you have completed this tutorial, you can go to other Amazon Redshift resources to learn more about the concepts introduced in this guide or you can reset your environment to the previous state. You might want to keep the sample cluster running if you intend to try tasks in other Amazon Redshift guides. However, remember that **you will continue to be charged for your cluster as long as it is**

running. You should revoke access to the cluster and delete it when you no longer need it so that you stop incurring charges.

Where Do I Go From Here?

Additional Resources

We recommend that you continue to learn more about the concepts introduced in this guide with the following resources:

- [Amazon Redshift Management Overview](#): This topic provides an overview of Amazon Redshift.
- [Amazon Redshift Cluster Management Guide](#): This guide builds upon this *Amazon Redshift Getting Started* and provides in-depth information about the concepts and tasks for creating, managing, and monitoring clusters.
- [Amazon Redshift Database Developer Guide](#): This guide builds upon this *Amazon Redshift Getting Started* by providing in-depth information for database developers about designing, building, querying, and maintaining the databases that make up your data warehouse.

Resetting Your Environment

When you have completed this tutorial, you should reset your environment to the previous state by doing the following:

- Revoke access to the port and CIDR/IP address for which you authorized access:


If you used the EC2-VPC platform to launch your cluster, perform the steps in [To Revoke Access from the VPC Security Group](#) (p. 18).

If you used the EC2-Classic platform to launch your cluster, perform the steps in [To Revoke Access from the Cluster Security Group](#) (p. 19).

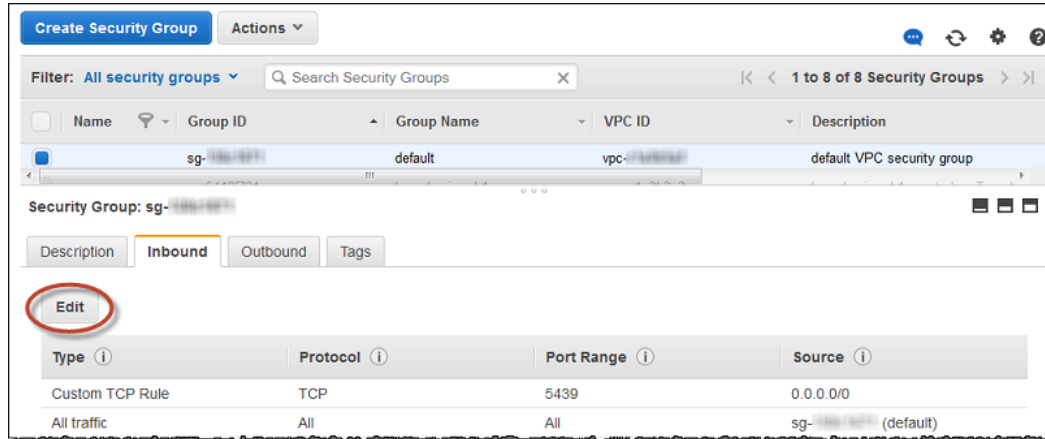
- Delete your sample cluster. **You will continue to incur charges for the Amazon Redshift service until you delete the cluster.** Perform the steps in [To Delete the Sample Cluster](#) (p. 20).

To Revoke Access from the VPC Security Group


1. In the Amazon Redshift console, in the navigation pane, choose **Clusters**.
2. Choose **examplecluster** to open it, and make sure you are on the **Configuration** tab.
3. Under **Cluster Properties**, choose the VPC security group.



4. With the default security group selected, choose the **Inbound** tab and then choose **Edit**.




5. Delete the custom TCP/IP ingress rule that you created for your port and CIDR/IP address 0.0.0.0/0. Do not remove any other rules, such as the **All traffic** rule that was created for the security group by default. Choose **Save**.




To Revoke Access from the Cluster Security Group

1. In the Amazon Redshift console, in the navigation pane, choose **Clusters**.
2. Choose examplecluster to open it, and make sure you are on the **Configuration** tab.
3. Under **Cluster Properties**, for **Cluster Security Groups**, choose **default** to open the default security group.



4. On the **Security Groups** tab, in the cluster security group list, choose the default cluster security group.
5. On the **Security Group Connections** tab, select the custom CIDR/IP ingress rule that you created for CIDR/IP address 0.0.0.0/0 and choose **Revoke**.



To Delete the Sample Cluster

1. In the Amazon Redshift console, in the navigation pane, choose **Clusters**.
2. Choose examplecluster to open it, and make sure you are on the **Configuration** tab.
3. In the **Cluster** menu, choose **Delete**.



4. In the **Delete Cluster** window, for **Create snapshot**, choose **No** and then choose **Delete**.

Delete Cluster [X]

If you want to restore your cluster later, you must create a manual snapshot. You can restore the cluster from the snapshot later and resume querying in minutes. You will be charged the Amazon S3 storage rate for a manual snapshot only if you don't have other clusters running, or if you exceed the allocated storage for active clusters.


Cluster: examplecluster ⓘ

Create snapshot: ☒ Yes ☐ No ⓘ

Snapshot name: ⓘ

[Cancel](#) [Delete](#)

5. On the cluster details window, the **Cluster Status** will display that the cluster being deleted.



Document History

The following table describes the important changes since the last release of the *Amazon Redshift Getting Started Guide*.

Latest documentation update: August 10, 2018

Change	Description	Release Date
New Feature	Updated the guide to describe the Quick launch cluster procedure.	August 10, 2018
New Feature	Updated the guide to launch clusters from the Amazon Redshift Dashboard.	July 28, 2015
New Feature	Updated the guide to use new node type names.	June 9, 2015
Documentation Update	Updated screenshots and procedure for configuring VPC security groups.	April 30, 2015
Documentation Update	Updated screenshots and procedures to match the current console.	November 12, 2014
Documentation Update	Moved loading data from Amazon S3 information into its own section and moved next steps section into the final step for better discoverability.	May 13, 2014
Documentation Update	Removed the Welcome page and incorporated the content into the main Getting Started page.	March 14, 2014
Documentation Update	This is a new release of the <i>Amazon Redshift Getting Started Guide</i> that addresses customer feedback and service updates.	March 14, 2014
New Guide	This is the first release of the <i>Amazon Redshift Getting Started Guide</i> .	February 14, 2013